

From Excel to Pandas: Advanced Data Analysis Techniques

Alexander Joel Molinar

October 16, 2024

1 From Excel to Pandas: Advanced Data Analysis Techniques

Excel is a powerful tool for data analysis, but when it comes to handling large datasets and performing complex operations, Python's pandas library offers superior capabilities. In this notebook, we'll explore advanced Excel features and demonstrate their pandas equivalents, focusing on pivot tables, charts, groupbys, and joins.

1.1 Generating a Random Sales Dataset

We'll create a random dataset to simulate sales data, which we'll use to demonstrate the advanced features.

```
[1]: import pandas as pd
import numpy as np
from faker import Faker
import random
import matplotlib.pyplot as plt
```

```
pd.set_option("display.notebook_repr_html", False)
```

```
[2]: # Initialize Faker
fake = Faker()

# Set the random seed for reproducibility
Faker.seed(0)
np.random.seed(0)
random.seed(0)

# Generate random data
num_records = 100000
data = {
    "Order ID": np.arange(1, num_records + 1),
    "Order Date": [fake.date_between(start_date='-1y', end_date='today') for _ in
    ↪in range(num_records)],
    "Customer Name": [fake.name() for _ in range(num_records)],
    "Segment": np.random.choice(["Consumer", "Corporate", "Home Office"],
    ↪num_records),
```

```

    "Category": np.random.choice(["Furniture", "Office Supplies",
↪ "Technology"], num_records),
    "Sub-Category": np.random.choice(["Bookcases", "Chairs", "Tables",
↪ "Accessories", "Appliances", "Art", "Binders", "Envelopes", "Fasteners",
↪ "Labels", "Paper", "Storage", "Supplies", "Phones", "Copiers", "Machines"],
↪ num_records),
    "Sales": np.round(np.random.uniform(10.0, 1000.0, num_records), 2),
    "Quantity": np.random.randint(1, 10, num_records),
    "Discount": np.round(np.random.uniform(0.0, 0.3, num_records), 2),
    "Profit": np.round(np.random.uniform(-200.0, 200.0, num_records), 2)
}

# Create a DataFrame
df = pd.DataFrame(data)

# Convert 'Order Date' to datetime
df['Order Date'] = pd.to_datetime(df['Order Date'])

# Preview the DataFrame
display(df.head())

```

	Order ID	Order Date	Customer Name	Segment	Category \
0	1	2024-09-06	Jeffrey Torres	Consumer	Technology
1	2	2024-03-12	Mrs. Jennifer King	Corporate	Furniture
2	3	2024-08-04	Kimberly Payne	Consumer	Office Supplies
3	4	2024-09-24	Paula Brady	Corporate	Furniture
4	5	2024-03-26	Jessica Archer	Corporate	Office Supplies

	Sub-Category	Sales	Quantity	Discount	Profit
0	Tables	130.30	9	0.18	84.39
1	Supplies	196.84	4	0.19	100.31
2	Fasteners	118.06	2	0.12	102.62
3	Fasteners	355.70	7	0.07	-58.91
4	Art	898.85	6	0.05	169.42

1.2 1. Pivot Tables

1.2.1 Excel Equivalent

In Excel, you can create a pivot table to summarize sales by category and segment.

1.2.2 Pandas Equivalent

```

[3]: # Updated code without the FutureWarning
pivot_table = pd.pivot_table(
    df,
    values='Sales',
    index=['Category'],

```

```

columns=['Segment'],
aggfunc='sum',
margins=True,
margins_name='Total'
)

print(pivot_table)

```

Segment	Consumer	Corporate	Home Office	Total
Furniture	5594767.38	5537224.51	5606353.76	16738345.65
Office Supplies	5574308.71	5628310.98	5602843.59	16805463.28
Technology	5625510.40	5628158.84	5573417.03	16827086.27
Total	16794586.49	16793694.33	16782614.38	50370895.20

1.3 2. GroupBy Operations

1.3.1 Excel Equivalent

Using SUMIF or array formulas to calculate total profit by category.

1.3.2 Pandas Equivalent

```

[4]: # Group by 'Category' and calculate total profit
category_profit = df.groupby('Category')['Profit'].sum().reset_index()
category_profit

```

```

[4]:
   Category  Profit
0  Furniture  27327.19
1  Office Supplies -13836.24
2  Technology  14366.78

```

1.4 3. Merging and Joining Data

1.4.1 Excel Equivalent

Using VLOOKUP or INDEX-MATCH to combine data from different sheets.

1.4.2 Pandas Equivalent

Let's assume we have another DataFrame with customer information.

```

[5]: # Create a customer DataFrame
customer_data = {
    "Customer Name": df['Customer Name'].unique(),
    "Region": np.random.choice(["East", "West", "Central", "South"],
    len(df['Customer Name'].unique()),
    "Customer Segment": np.random.choice(["High Value", "Medium Value", "Low
    Value"], len(df['Customer Name'].unique()))
}

```

```
customers = pd.DataFrame(customer_data)
customers.head()
```

```
[5]:      Customer Name Region Customer Segment
0      Jeffrey Torres  South      High Value
1  Mrs. Jennifer King   West      Low Value
2      Kimberly Payne   West      High Value
3        Paula Brady    East      Low Value
4      Jessica Archer   East      Low Value
```

```
[6]: # Merge the sales data with customer data
merged_df = pd.merge(df, customers, on='Customer Name', how='left')
merged_df.head()
```

```
[6]:      Order ID Order Date      Customer Name      Segment      Category \
0          1 2024-09-06      Jeffrey Torres      Consumer      Technology
1          2 2024-03-12  Mrs. Jennifer King      Corporate      Furniture
2          3 2024-08-04      Kimberly Payne      Consumer      Office Supplies
3          4 2024-09-24        Paula Brady      Corporate      Furniture
4          5 2024-03-26      Jessica Archer      Corporate      Office Supplies

      Sub-Category      Sales      Quantity      Discount      Profit Region Customer Segment
0          Tables      130.30           9          0.18      84.39  South      High Value
1          Supplies      196.84           4          0.19     100.31   West      Low Value
2          Fasteners      118.06           2          0.12     102.62   West      High Value
3          Fasteners      355.70           7          0.07     -58.91   East      Low Value
4           Art      898.85           6          0.05     169.42   East      Low Value
```

1.5 4. Data Visualization (Charts)

1.5.1 Excel Equivalent

Creating charts to visualize sales trends over time.


1.5.2 Pandas/Matplotlib Equivalent

```
[7]: # Calculate monthly sales
monthly_sales = df.set_index('Order Date').resample('M')['Sales'].sum()

# Plot the monthly sales
plt.figure(figsize=(12,6))
plt.plot(monthly_sales.index, monthly_sales.values, marker='o')
plt.title('Monthly Sales Over the Past Year')
plt.xlabel('Month')
plt.ylabel('Sales')
plt.grid(True)
plt.show()
```

C:\Users\AlexanderMolinar\AppData\Local\Temp\ipykernel_41872\3861109624.py:2:
FutureWarning: 'M' is deprecated and will be removed in a future version, please
use 'ME' instead.

```
monthly_sales = df.set_index('Order Date').resample('M')['Sales'].sum()
```



1.6 5. Advanced GroupBy with Multiple Aggregations

1.6.1 Excel Equivalent

Using Pivot Tables with multiple value fields and custom calculations.

1.6.2 Pandas Equivalent

```
[8]: # Group by 'Category' and 'Sub-Category' and calculate total Sales and Profit
grouped = df.groupby(['Category', 'Sub-Category']).agg({'Sales': 'sum',
↪ 'Profit': 'sum'}).reset_index()
grouped.head()
```

```
[8]:   Category Sub-Category   Sales  Profit
0  Furniture  Accessories 1058604.06  4045.37
1  Furniture   Appliances 1054785.39  4191.92
2  Furniture         Art 1034452.53 -8202.14
3  Furniture   Binders 1037405.95  3796.30
4  Furniture   Bookcases 1047625.21 -1685.37
```

1.7 6. Applying Custom Functions with GroupBy

1.7.1 Excel Equivalent

Using array formulas or VBA scripts to apply custom calculations.

1.7.2 Pandas Equivalent

```
[9]: # Define a custom function to calculate profit margin
def profit_margin(x):
    return (x['Profit'].sum() / x['Sales'].sum()) * 100

# Apply the custom function
margin = df.groupby('Category').apply(profit_margin).reset_index(name='Profit_
↳Margin (%)')
margin
```

C:\Users\AlexanderMolinar\AppData\Local\Temp\ipykernel_41872\1944146675.py:6:
DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns.
This behavior is deprecated, and in a future version of pandas the grouping
columns will be excluded from the operation. Either pass `include_groups=False`
to exclude the groupings or explicitly select the grouping columns after groupby
to silence this warning.

```
margin = df.groupby('Category').apply(profit_margin).reset_index(name='Profit
Margin (%)')
```

```
[9]:      Category  Profit Margin (%)
0      Furniture      0.163261
1  Office Supplies     -0.082332
2      Technology      0.085379
```

1.8 7. Filtering Data

1.8.1 Excel Equivalent

Using AutoFilter or advanced filtering options.

1.8.2 Pandas Equivalent

```
[10]: # Filter orders with Sales > $500 and Profit < 0
filtered_df = df[(df['Sales'] > 500) & (df['Profit'] < 0)]
filtered_df.head()
```

```
[10]:   Order ID Order Date Customer Name Segment Category \
6      7  2024-01-23 Margaret Valentine Consumer Office Supplies
11     12  2024-08-14 Amy Cunningham Home Office Technology
12     13  2024-09-01 Donald Miller Corporate Technology
16     17  2024-05-28 John Zamora Corporate Office Supplies
21     22  2023-12-08 Jared Alexander Corporate Furniture
```

	Sub-Category	Sales	Quantity	Discount	Profit
6	Labels	743.11	8	0.22	-137.31
11	Appliances	574.29	6	0.26	-34.20
12	Envelopes	796.95	1	0.22	-115.04
16	Binders	848.17	8	0.16	-129.14
21	Supplies	912.02	4	0.16	-23.38

1.9 8. Handling Missing Data

1.9.1 Excel Equivalent

Using ISBLANK or filtering out empty cells.

1.9.2 Pandas Equivalent

Let's introduce some missing data into our DataFrame.

```
[11]: # Introduce missing values
df_missing = df.copy()
df_missing.loc[df_missing.sample(frac=0.05).index, 'Profit'] = np.nan

# Check for missing values
df_missing['Profit'].isnull().sum()
```

```
[11]: 5000
```

```
[12]: # Fill missing Profit values with 0
df_missing['Profit'].fillna(0, inplace=True)

# Verify missing values are handled
df_missing['Profit'].isnull().sum()
```

C:\Users\AlexanderMolinar\AppData\Local\Temp\ipykernel_41872\3140221436.py:2:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using
'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)
instead, to perform the operation inplace on the original object.

```
df_missing['Profit'].fillna(0, inplace=True)
```

```
[12]: 0
```

1.10 9. Sorting Data

1.10.1 Excel Equivalent

Using the Sort function to organize data.

1.10.2 Pandas Equivalent

```
[13]: # Sort data by Sales in descending order
sorted_df = df.sort_values(by='Sales', ascending=False)
sorted_df.head()
```

```
[13]:      Order ID Order Date  Customer Name  Segment  Category \
71151      71152 2024-10-04    Steve Wolfe  Consumer  Office Supplies
71815      71816 2023-12-14    Chad Williams  Corporate      Furniture
41690      41691 2024-06-10  Jonathan Keller  Home Office      Technology
66577      66578 2023-11-17    Jeffrey Lynch  Corporate  Office Supplies
32484      32485 2023-12-21      Lisa Jones  Corporate  Office Supplies
```

	Sub-Category	Sales	Quantity	Discount	Profit
71151	Labels	1000.00	6	0.27	-25.93
71815	Appliances	999.98	3	0.27	180.22
41690	Bookcases	999.96	2	0.21	28.51
66577	Paper	999.96	4	0.18	3.54
32484	Bookcases	999.94	9	0.05	-126.83

1.11 10. Conditional Formatting

1.11.1 Excel Equivalent

Using Conditional Formatting to highlight cells based on values.

1.11.2 Pandas Equivalent

```
[14]: # Define a function to highlight negative profits
def highlight_negative(s):
    return ['background-color: red' if v < 0 else '' for v in s]

# Apply the function to the DataFrame
styled_df = df[['Order ID', 'Sales', 'Profit']].head(10).style.
    ↪ apply(highlight_negative, subset=['Profit'])
styled_df
```

```
[14]: <pandas.io.formats.style.Styler at 0x26bdc55ce30>
```

1.12 Conclusion

Pandas offers a rich set of functionalities that can replicate and often enhance the capabilities of Excel's advanced features. Transitioning to pandas allows for more efficient data manipulation,

especially with large datasets. By understanding these pandas equivalents, you can take your data analysis skills to the next level.

Feel free to use this notebook as a guide for advanced data analysis techniques in pandas. The combination of pandas and Python provides a powerful platform for handling, analyzing, and visualizing data.